Online Variants of Value Iteration and Approximation Via State Aggregation

Daniel Russo

April 13, 2020

Columbia University

References

These slides try to synthesize material from many papers.

Real time value iteration and optimistic exploration

- What I call "Real-time value iteration" (RTVI) is typically called 'Real time dynamic programming, due to Barto et al. [1995].
    - The regret style analysis here appears to be new. Asymptotic convergence is due to Barto et al. [1995].
- While we won't cover it, the foundations of many optimistic exploration algorithms in RL can be understood through the lens of RTVI. Some important related papers include Kearns and Singh [2002], Brafman and Tennenholtz [2002], Strehl et al. [2006], Jaksch et al. [2010], Azar et al. [2017], Jin et al. [2018].

References

State Aggregation

- Work by Gordon [1995] and Tsitsiklis and Van Roy [1996] explains that fitted value iteration converges tp a fixed point with state-aggregation, but not in general.
- Van Roy [2006] explains the crucial performance gains due to finding a fixed point in the right state-relevance distribution.
- Li et al. [2006] formalizes the notion that our assumptions that with regard to state-aggregation are weaker than e.g. assuming the model dynamics are sufficiently smooth.
- These slides are inspired by the preprint Dong et al. [2019].

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Where are we heading?
Value function approximation

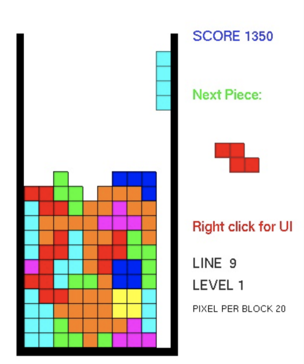

- State $s \in \{0,1\}^{10\times 20}$ is a board configuration.
- One approach is to fit a linear approximation: $J^*(s) \approx J_\theta(s) := \phi(s)^\top \theta$.
- $\phi(s)$ encodes features. E.g. column heights, inter-column height differences, max height etc.
- Greedy action under J_θ is simple: move the piece to the position with minimal estimated cost-to-go.

Fitted Value Iteration (FVI)

Regression based approximation to Bellman updates.

- Define the weighted norm $\|J\|_{2,\nu} = \sqrt{\mathbb{E}_{s\sim\nu}[J(s)^2]}$.
- Fitted value iteration is the scheme: for $k = 1, 2, \cdots$

$$\theta_{k+1} \in \operatorname*{argmin}_{\theta\in\Theta} \|J_\theta - TJ_{\theta_k}\|_{2,\nu}$$

- Equivalently, this can be viewed as a projected value iteration:

$$J_{k+1} = \Pi_{\mathcal{F},\nu} TJ_k$$

where $\mathcal{F} = \{J_\theta \mid \theta \in \Theta\}$ is the space of value functions approximations and $\Pi_{\mathcal{F},\nu} J = \operatorname{argmin}_{f\in\mathcal{F}} \|f - J\|_{2,\nu}$ projects onto $\mathcal{F}$ in a weighted norm.

Sample based approximations to FVI

In practice, we typically approximate expectations by random sampling.

1. Sample n states $s_1, \cdots, s_n \sim \rho(\cdot)$.
2. For each state s_i, and each control $u \in U(s)$, draw m samples of the successor states $s_{i,u}^{(1)}, \cdots, s_{i,u}^{(m)} \sim P_{s_i,\cdot}(u)$,
3. For $k = 0, 1, 2, \cdot$

$$\theta_{k+1} \in \operatorname*{argmin}_{\theta} \frac{1}{n} \sum_{i=1}^{n} \left(J_\theta(s_i) - \underbrace{\left(\min_u g(s_i, u) + \gamma \frac{1}{m} \sum_{j=1}^{m} J_{\theta_k}(s_{i,u}^{(j)}) \right)}_{\text{sample approx. to } TJ_{\theta_k}(s_i)} \right)$$

Comments:

(a) If state transitions are very sparse, $TJ(s)$ can be computed exactly.

(b) Statistical learning theory bounds the error from random sampling.

Fitted Q-iteration

It is simpler to approximate Bellman updates to Q-functions.

- Define the state-action value function:

$$Q^*(s,u) = g(s,u) + \alpha \sum_{s' \in S} P_{ss'}(u) J^*(s')$$

- Obeys the Bellman $Q^* = FQ^*$ where

$$FQ(s,u) := g(s,u) + \alpha \sum_{s' \in S} P_{ss'}(u) \min_{u'} Q(s', u')$$

- Fitted Q iteration is the scheme: for $k = 1, 2, \cdots$

$$\theta_{k+1} \in \operatorname*{argmin}_{\theta \in \Theta} \| Q_\theta - FQ_{\theta_k} \|_{2,\nu} \tag{1}$$

  where ν is a distribution over state, control pairs.

The Q-learning algorithm essentially makes a stochastic gradient updates rather than solving (1) *exactly.*

Does this work?

1. Convergence to a fixed point $\widehat{J} = \Pi_{\rho,\mathcal{F}} T \widehat{J}$?
2. Does it produce an accurate approximation if J^* is “close to” the function class $\mathcal{F}$?
3. Is the resulting policy near optimal?
4. How should we set the state-importance-weights ν?

Unfortunately . . . this procedure does not converge in general. Existing guarantees on typically performance require extremely strong assumptions.

Plan for today

1. State-aggregation: a simple case of value function approximation with which fitted-value-iteration is convergent.

2. Understanding the state-importance-weights and how this should be adapted over time.

State-Aggregation (a.k.a state abstraction)

- We believe $J^*(s) \approx J^*(\tilde{s})$ if s and $\tilde{s}$ are "similar."

- More formally, take $\phi : \mathcal{S} \to \mathcal{S}$ to a be a mapping that associates $s \in \mathcal{S}$ with a representative state $\phi(s) \in \mathcal{S}$.
  - Simple case is $\mathcal{S} = [0, 1]$ and ϕ maps $s \in [0, .01)$ to $\phi(s) = .005$, $s \in [.01, .02)$ to $\phi(s) = .015$ and so on.

- State aggregated value functions are

$$\begin{aligned} \mathcal{F}_\phi &= \{f | f(s) = f(\phi(s)) \, \forall s\} \\ &= \{f | f(s) = f(\tilde{s}) \quad \text{if } \phi(s) = \phi(\tilde{s})\} \end{aligned}$$

  - In the simple case described above, $J \in \mathcal{F}_\phi$ is defined by the 99 values $J(.005), J(.015), \cdots, J(.995)$.

- $\Phi = \{\phi(s) : s \in \mathcal{S}\} = \{1, \cdots, m\}$. (. . . w.l.o.g)
  Set $\mathcal{S}_i = \{s \in \mathcal{S} : \phi(s) = i\}$ for $i = 1, \cdots, m$.

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Convergence of FVI with state-aggregation

Theorem
Consider the iteration

$$J_{k+1} = \Pi_{\mathcal{F}_\phi,\nu} T J_k \quad k = 0, 1, \cdots$$

where $\nu(s) > 0$ for all $s \in \mathcal{S}$. Then,

$$\|J_k - \widehat{J}\|_\infty \leq \alpha^k \|J_0 - \widehat{J}\|_\infty$$

where $\widehat{J}$ solves the projected Bellman equation

$$\widehat{J} = \Pi_{\mathcal{F}_\phi,\nu} T \widehat{J}.$$

Convergence proof

Crucial fact: $\Pi_{\mathcal{F}_\phi,\nu}$ is a non-expansion in $\|\cdot\|_\infty$.

Proof.

$$\Pi_{\mathcal{F}_\phi,\nu} J \in \operatorname*{argmin}_{\widehat{J}\in\mathcal{F}_\phi} \mathbb{E}_{s\sim\nu}\left[\left(\widehat{J}(s) - J(s)\right)^2\right]$$

It is solved by the conditional mean

$$\widehat{J}(i) = \mathbb{E}_{s\sim\nu}\left[J(s) \mid s \in \mathcal{S}_i\right]$$

□

Another fact: $\Pi_{\mathcal{F}_\phi,\nu}$ is linear.

Convergence proof

We show the projected Bellman update is a max-norm contraction. The convergence result follows immediately.

<u>**Lemma**</u> $\Pi_{\mathcal{F}_\phi,\nu}T$ is a contraction w.r.t $\|\cdot\|_\infty$ with modulus of contraction α.

Proof.

$$\begin{aligned} \|\Pi_{\mathcal{F}_\phi,\nu}TJ - \Pi_{\mathcal{F}_\phi,\nu}T\bar{J}\|_\infty &= \|\Pi_{\mathcal{F}_\phi,\nu}\left(TJ - T\bar{J}\right)\|_\infty \\ &\leq \|TJ - T\bar{J}\|_\infty \\ &\leq \alpha\|J - \bar{J}\|_\infty \end{aligned}$$

□

Approximation error bound

Is $\hat{J}$ an effective approximation to J^*? We compare its quality to the best possible approximation possible with state-aggregation.

Theorem
Set $\epsilon = \|\Pi_{\mathcal{F}_\phi,\nu} J^* - J^*\|_\infty$. *Then*

$$\|\widehat{J} - J^*\|_\infty \leq \frac{\epsilon}{1-\alpha}$$

Proof.

$$\begin{aligned} \|\widehat{J} - J^*\|_\infty &\leq \|\widehat{J} - \Pi_{\mathcal{F}_\phi,\nu} J^*\|_\infty + \|\Pi_{\mathcal{F}_\phi,\nu} J^* - J^*\|_\infty \\ &= \|\Pi_{\mathcal{F}_\phi,\nu} T\widehat{J} - \Pi_{\mathcal{F}_\phi,\nu} TJ^*\|_\infty + \epsilon \\ &\leq \alpha\|\widehat{J} - J^*\|_\infty + \epsilon. \end{aligned}$$

□

<u>**Comment**</u>: We can't directly compute the projection of J^*, since we don't know it. We instead solve a projected version of Bellman's equation, which leads the error bound to expand by $1/(1-\alpha)$.

Performance loss bound

How effective is the greedy policy computed with respect to $\widehat{J}$?

Theorem
Let $\mu \in G(\widehat{J})$ *and* $\epsilon = \|\Pi_{\mathcal{F}_\phi,\nu} J^* - J^*\|_\infty$. *Then,*

$$\|J_\mu - J^*\|_\infty \leq \frac{\alpha\|\widehat{J} - J\|_\infty}{1-\alpha} \leq \frac{\alpha\epsilon}{(1-\alpha)^2}$$

Proof.
The first inequality was shown last class. □

Tightness of the Approximation Error Bound

The dependence on $1/(1-\alpha)^2$ is highly problematic. Unfortunately, it is tight due to an example discussed in detail in Van Roy (2006).

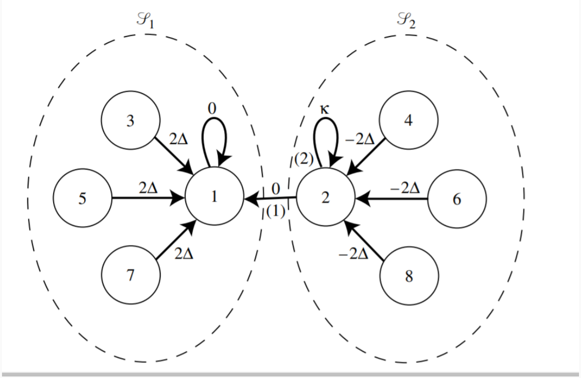

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Reducing performance loss via state-relevance weighting

Recall the discounted state occupancy measure

$$d_{\infty}^{\mu} = (1-\alpha)\sum_{t=0}^{\infty} \alpha^{t} d_0 P_{\mu}^{t}.$$

We will show a natural online algorithm converges toward a fixed point

$$\widehat{J} = \Pi_{\mathcal{F}_{\phi}, d_{\infty}^{\mu}} T\widehat{J} \qquad \mu \in G(\widehat{J}).$$

This satisfies the performance loss bound

$$\mathbb{E}_{s\sim d_0}\left[J_{\mu}(s) - J^{*}(s)\right] \leq \frac{\epsilon}{1-\alpha}$$

Comments:

- Saves a factor of $1/(1-\alpha)$ in the worst-case.
- The state-relevance-weighting is the fraction of time spent in a given state under the selected policy.
- This is a fixed point in $\{$the space of cost-to-functions$\} \times \{$the space of policies$\}$.

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Real time value iteration

Real Time Value Iteration

- Input J_0
- For episode $k = 0, 1, \cdots$
    1. Select greedy policy $\mu_k \in G(J_k)$
    2. Draw random initial state $s_0^{(k)} \sim d_0$.
    3. Sample episode length $\tau \sim \mathrm{Geom}(1-\alpha)$
    4. Sample $(s_0^{(k)}, \cdots, s_\tau^{(k)})$ by applying μ_k for τ timesteps.
    5. Make Bellman update at $s_\tau^{(k)}$ (. . . Or at all visited states)

$$J_{k+1}(s) \leftarrow \begin{cases} TJ_k(s) & \text{if } s = s_\tau^{(k)} \\ J_k(s) & \text{if } s \neq s_\tau^{(k)} \end{cases}$$

Real time value iteration with Q functions

- Closer to what is used in RL, since no knowledge of the environment is required to compute a greedy policy w.r.t Q.

Real Time Value Iteration with Q functions

- Input Q_0
- For episode $k = 0, 1, \cdots$
  1. Select greedy policy $\mu_k \in G(Q_k)$
  2. Draw random initial state $s_0^{(k)} \sim d_0$.
  3. Sample episode length $\tau \sim \mathrm{Geom}(1-\alpha)$
  4. Sample $(s_0^{(k)}, \cdots, s_\tau^{(k)})$ by applying μ_k for τ timesteps.
  5. Make Bellman update at $s_\tau^{(k)}$ (...Or at all visited states)

$$Q_{k+1}(s,u) \leftarrow \begin{cases} FQ_k(s,u) & \text{if } s = s_\tau^{(k)}, u = \mu_k(s_\tau^{(k)}) \\ Q_k(s,u) & \text{if otherwise} \end{cases}$$

Q-learning

Dropping the k superscripts for ease of notation

Q-learning with greedy exploration

- Input initial Q
- For episode $k = 0, 1, \cdots$
  1. Select greedy policy $\mu \in G(Q)$
  2. Draw random initial state $s_0 \sim d_0$.
  3. Sample episode length $\tau \sim \mathrm{Geom}(1-\alpha)$
  4. Sample $(s_0, \cdots, s_\tau, s_{\tau+1})$ by applying μ for τ timesteps.
  5. Define $u_t = \mu(s_t)$ for $t = 0, \cdots, \tau + 1$.
  6. Make soft noisy Bellman update at s_τ (. . . Or $s_0 \cdots s_\tau$)

$$Q(s_\tau, \mu(s_\tau)) \leftarrow (1-\beta_k) Q(s_\tau, \mu(s_\tau)) + \beta_k \underbrace{\Big[g(s_\tau, u_\tau) + \gamma \min_u Q(s_{\tau+1}, u)}_{\text{Unbiased obsrvation of } FQ(s_\tau, u_\tau)}$$

Convergence of Real-time Value Iteration?

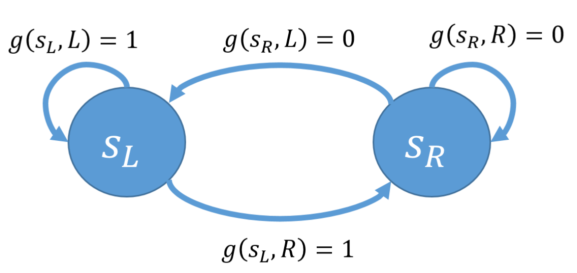

- Suppose the initial state is always L ($d_0 = (1, 0)$).
- Suppose $J = (1, 2)/(1 - \alpha)$,
  - Greedy policy $\mu \in G(J)$ satisfies $\mu(s_L) = L$.
  - The system always stays in state L
  - Also $TJ(s_L) = J(s_L)$.
- Suppose $J(s_R) \leq 0$. [. . . *Optimism in the face of uncertainty*]
  - If $G(J) = (R, R)$ then the induced policy is optimal.
  - Otherwise, we must have $J(s_L) < J^*(s_L)$.
    Real-time VI increases the estimate $J(s_L)$.

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Error bounds depend on the state distribution

Notation: Define $J(d) = \sum_{s \in \mathcal{S}} d(s) J(s)$.

Lemma from last class. For any $J \in \mathbb{R}^n$ and policy μ',

$$J - J_{\mu'} = (I - \alpha P_{\mu'})^{-1} (J - T_{\mu'} J)$$

$$J(d_0) - J_{\mu'}(d_0) = \frac{1}{1-\alpha} (J - T_{\mu'} J)(d_\infty^{\mu'})$$

Performance loss bounds that depend on the state distribution

Lemma: For any $J \in \mathbb{R}^n, \mu \in G(J)$ and optimal policy μ^*,

$$
\begin{aligned}
0 \preceq \quad & J_\mu - J^* \preceq \left[(I - \alpha P_{\mu^*})^{-1} - (I - \alpha P_\mu)^{-1}\right](J - TJ) \\
0 \preceq \quad & J_\mu(d_0) - J^*(d_0) \preceq (J - TJ)(d_\infty^{\mu^*}) - (J - TJ)(d_\infty^{\mu})
\end{aligned}
$$

Proof.
Applying the earlier Lemma with $\mu' = \mu^*$ gives

$$J - J^* = (I - \alpha P_{\mu^*})^{-1}(J - T_{\mu^*}J) \preceq (I - \alpha P_{\mu^*})^{-1}(J - TJ)$$

Since $\mu \in G(J)$, we also have

$$J - J_\mu = (I - \alpha P_\mu)^{-1}(J - T_\mu J) = (I - \alpha P_\mu)^{-1}(J - TJ)$$

Subtracting yields the result. □

Comment: For our current purposes, the dependence on μ^* is problematic, as it's unknown.

Optimism to the rescue

Performance loss is bounded by *on policy* Bellman errors.

Lemma: If $J \preceq J^*$, $\mu \in G(J)$,

$$0 \preceq J_\mu - J^* \preceq (I - \alpha P_\mu)^{-1}(TJ - J)$$

$$0 \preceq J_\mu(d_0) - J^*(d_0) \preceq (TJ - J)(d_\infty^\mu)$$

Proof.
$J_\mu - J^* \preceq J_\mu - J = (I - \alpha P_\mu)^{-1}(TJ - J)$. □

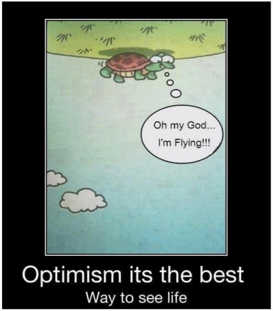

Regret Style Analysis of RTVI

Assumption: Optimistic initiation $J_0 = \frac{-M}{1-\alpha}$. (Recall $M = \|g\|_\infty$)

Keys to analysis:

- **On policy sampling**: The state $s_\tau^{(k)}$ is sampled from $d_\infty^{\mu_k}$.

$$\begin{aligned}\mathbb{P}(s_\tau^{(k)} = s) = \mathbb{E}\left[\mathbb{P}(s_\tau^{(k)} = s \mid \tau)\right] &= \mathbb{E}\left[(d_0 P_{\mu_k}^\tau)(s)\right] \\ &= (1-\alpha)\sum_{t=0}^{\infty} \alpha^t (d_0 P_{\mu_k}^t(s)).\end{aligned}$$

- **Monotonicity**: Since $J_0 \preceq TJ_0$, we have

$$J_0 \preceq TJ_0 \preceq T^2 J_0 \preceq \cdots \preceq J^*$$

and one can similarly show the iterates of RTVI are monotone:

$$J_0 \preceq J_1 \preceq J_2 \cdots$$

Regret Style Analysis of RTVI

Set $\mathrm{Sum}(J) = \sum_{s \in \mathcal{S}} J(s)$. By our earlier Lemma,

$$
\begin{aligned}
(1-\alpha)(J_{\mu_0}(d_0) - J^*(d_0)) \leq (TJ_0 - J_0)(d_\infty^{\mu_0}) &= \mathbb{E}\left[(TJ_0 - J_0)(s_\tau^{(0)})\right] \\
&= \mathbb{E}\left[(J_1 - J_0)(s_\tau^{(0)})\right] \\
&= \mathbb{E}\left[\mathrm{Sum}(J_1) - \mathrm{Sum}(J_0)\right]
\end{aligned}
$$

Applying this for each episode k, and using the tower property:

$$
\begin{aligned}
\mathrm{Regret(K)} &:= \mathbb{E}\left[\sum_{k=1}^{K} J_k(d_0) - J^*(d_0)\right] \\
&\leq (1-\alpha)^{-1}\mathbb{E}\left[\sum_{k=0}^{K} (\mathrm{Sum}(J_k) - \mathrm{Sum}(J_{K+1}))\right] \\
&= (1-\alpha)^{-1}\mathrm{Sum}(J_0) - \mathrm{Sum}(J_{k+1}) \\
&\leq (1-\alpha)^{-1}\mathrm{Sum}(J_0 - J^*) \\
&\leq \frac{2M|\mathcal{S}|}{(1-\alpha)^2}.
\end{aligned}
$$

Sketch of asymptotic analysis of RTVI

We converge to a cost-to-go function satisfying Bellman's equation at all states visited with positive probability.

Optimism implies optimality.

- Since $\{J_k\}$ is a monotone sequence, it has a limit $J_\infty \preceq J^*$.
- Let $\mu_\infty \in G(J_\infty)$ be the corresponding policy, and assume it's unique so $\mu_k = \mu_\infty$ for sufficiently large k.
- Since the limit exists, we must have

$$\|J_k - J_{k+1}\|_\infty \to 0$$

- This implies

$$J_k(s) - TJ_k(s) \to 0$$

for s visited with positive probability ($d_\infty^{\mu_\infty}(s) > 0$). So

$$J_{\mu_\infty}(d_0) - J^*(d_0) \leq (TJ_\infty - J_\infty)(d_\infty^{\mu_\infty}) = 0.$$

Table of Contents

Preview of value function approximation

Fitted value iteration with state aggregation

An online algorithm provides a better state-relevance weighting

Real time value iteration and optimistic exploration

Convergence analysis of real time value iteration

Real-time value iteration with state-aggregation

Which algorithm attains this performance bound?

We use require $\epsilon = \sup_s \|J^*(\phi(s)) - J^*(s)\|$.

Real Time Value Iteration with State Aggregation

- Input J_0 with $J_0(s) = \frac{-M}{1-\alpha}$ $\forall s$. (…Recall $M = \|g\|_\infty$)
- For $k = 0, 1, \cdots$
  1. Sample $s_k \sim d_\infty^{\mu_k}$ where $\mu_k \in G(J_k)$
  2. Make a conservative update to the representative state $\phi(s_k)$:

$$J_{k+1}(s) \leftarrow \begin{cases} TJ_k(s_k) - 2\epsilon & \text{if } \phi(s) = \phi(s_k) \\ J_k(s) & \text{if } \phi(s) \neq \phi(s_k) \end{cases}$$

Comments:

- Step 1 can be executed by simulating a state trajectory under a greedy policy with respect to J_k, as shown before.
- An efficient implementation only needs to store the value at the representative state.
- With Q-function variants it is easier to apply a greedy policy.

Regret style analysis of RTVI with state-aggregation

Notation: Representative states: $\Phi = \{\phi(s) : s \in \mathcal{S}\}$
$\mathrm{Sum}(J) := \sum_{s\in\Phi} J(s)$.

The next result establishes optimism for the iterates J_k of RTVI.
Lemma (Optimism): $J_k \preceq J^*$ for each k.

Proof.
We have $J_0 \preceq J^*$ by definition.
For every s with $\phi(s) = \phi(s_0)$ (. . . the states we update . . .),

$$J_1(s) = TJ_0(s_0) - 2\epsilon \leq TJ^*(s_0) - 2\epsilon \leq J^*(s)$$

where the inequality used that

$$TJ^*(s) = J^*(s) \leq J^*(s_0) + |J^*(s) - J^*(\phi(s_k)| + |J^*(s_0) - J^*(\phi(s_k))|.$$

□

Regret style analysis of RTVI with state-aggregation (continued)

As before, we have the bound in each episode:

$$\begin{aligned}(1-\alpha)\left[J_{\mu_0}(d_0)-J^*(d_0)\right] &\leq (TJ_0-J_0)\left(d_\infty^{\mu_0}\right)\\ &= \mathbb{E}\left[(TJ_0-J_0)(s_0)\right]\\ &\overset{(*)}{=} \mathbb{E}\left[(J_1-J_0)(\phi(s_0))\right]+2\epsilon\\ &= \mathbb{E}\left[\mathrm{Sum}(J_1)-\mathrm{Sum}(J_0)\right]+2\epsilon\end{aligned}$$

The equality (*) uses that by the definition of the algorithm

$$J_1(\phi(s_0)) = TJ_0(s_0)-2\epsilon.$$

Regret style analysis sof RTVI (continued)

Applying this for each episode k, and using the tower property:

$$\begin{aligned}
\mathrm{Regret(K)} &:= \mathbb{E}\left[\sum_{k=1}^{K} J_{\mu_k}(d_0) - J^*(d_0)\right] \\
&\leq \frac{1}{1-\alpha}\mathbb{E}\left[\sum_{k=0}^{K} (\mathrm{Sum}(J_{k+1}) - \mathrm{Sum}(J_K))\right] + \frac{2K\epsilon}{1-\alpha} \\
&= \frac{\mathrm{Sum}(J_{k+1}) - \mathrm{Sum}(J_0)}{1-\alpha} + \frac{2K\epsilon}{1-\alpha} \\
&\leq \frac{\mathrm{Sum}(J^* - J_0)}{1-\alpha} + \frac{2K\epsilon}{1-\alpha} \\
&\leq \frac{2M|\Phi|}{(1-\alpha)^2} + \frac{2K\epsilon}{1-\alpha}.
\end{aligned}$$

We see that average regret scales as

$$\limsup_{K\to\infty} \frac{\mathrm{Regret(K)}}{K} \leq \frac{2\epsilon}{1-\alpha}$$

Mohammad Gheshlaghi Azar, Ian Osband, and Rémi Munos. Minimax regret bounds for reinforcement learning. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 263–272. JMLR. org, 2017.

Andrew G Barto, Steven J Bradtke, and Satinder P Singh. Learning to act using real-time dynamic programming. *Artificial intelligence*, 72(1-2):81–138, 1995.

Ronen I Brafman and Moshe Tennenholtz. R-max-a general polynomial time algorithm for near-optimal reinforcement learning. *Journal of Machine Learning Research*, 3(Oct): 213–231, 2002.

Shi Dong, Benjamin Van Roy, and Zhengyuan Zhou. Provably efficient reinforcement learning with aggregated states. *arXiv preprint arXiv:1912.06366*, 2019.

Geoffrey J Gordon. Stable function approximation in dynamic programming. In *Machine Learning Proceedings 1995*, pages 261–268. Elsevier, 1995.

Thomas Jaksch, Ronald Ortner, and Peter Auer. Near-optimal regret bounds for reinforcement learning. *Journal of Machine Learning Research*, 11(Apr):1563–1600, 2010.

Chi Jin, Zeyuan Allen-Zhu, Sebastien Bubeck, and Michael I Jordan. Is q-learning provably efficient? In *Advances in Neural Information Processing Systems*, pages 4863–4873, 2018.

Michael Kearns and Satinder Singh. Near-optimal reinforcement learning in polynomial time. *Machine learning*, 49(2-3):209–232, 2002.

Lihong Li, Thomas J Walsh, and Michael L Littman. Towards a unified theory of state abstraction for mdps. In *In Proceedings of the Ninth International Symposium on Artificial Intelligence and Mathematics*. Citeseer, 2006.

Alexander L Strehl, Lihong Li, Eric Wiewiora, John Langford, and Michael L Littman. Pac model-free reinforcement learning. In *Proceedings of the 23rd international conference on Machine learning*, pages 881–888, 2006.

John N Tsitsiklis and Benjamin Van Roy. Feature-based methods for large scale dynamic programming. *Machine Learning*, 22 (1-3):59–94, 1996.

Benjamin Van Roy. Performance loss bounds for approximate value iteration with state aggregation. *Mathematics of Operations Research*, 31(2):234–244, 2006.